

• FREE GUIDE

AI Detection Field Guide 2026

How AI detectors work, why honest writing gets flagged, and how to write in a way that sounds like you. No hype, no tricks.

The short version

If you only read one page, read this one.

- 1. **Detectors estimate, they do not prove.** Every major tool gives a probability. Turnitin says its own score "should not be used as the sole basis" for action.
- 2. **Honest writing gets flagged.** In a Stanford study, seven detectors wrongly labeled over half of a set of essays by non-native English speakers as AI-generated.
- 3. **The triggers are mostly about predictability.** Stock phrases, even sentence rhythm, generic claims and a fixed template look like AI, whoever wrote them.
- 4. **The fix is better writing, not disguise.** Specific details, your own experience, varied rhythm and a real point of view make text more human, more readable and harder to confuse with a template.
- 5. **Nothing guarantees an outcome.** Detectors change often. Turnitin added detection of "AI bypasser" (humanizer) tools in August 2025. Anyone who promises "100% undetectable" is selling you a risk.

Keep a record of your writing process (drafts, version history, notes). If you are ever flagged, that record matters more than any score. Our free *False AI Flag Appeal Kit* covers what to do step by step.

How AI detectors work

Three families of detectors

Approach	What it looks at	Who uses it	Weak spot
Trained classifier	A machine-learning model trained on large sets of human and AI text learns which features separate them.	Turnitin, GPTZero (since 2023) and many other commercial tools.	It is only as good as its training data. Writing styles it saw less often (second-language writing, very formal or very simple prose) can be misjudged.
Statistical signals	How predictable each word is (perplexity) and how much that varies across the text (burstiness).	Early GPTZero; still used as one signal inside newer tools.	Plain, careful, predictable human writing scores like AI.
Watermarks	A hidden statistical pattern that the AI model itself adds while generating text, which only a matching detector can read.	Google's SynthID Text, used in Gemini and released as open source in October 2024.	Only works on text from models that add the watermark. Google says it is less reliable on short text, on thoroughly rewritten text and on translated text.

What Turnitin says about its own method

In its 2023 FAQ, Turnitin explained that it splits a submission into overlapping segments of a few hundred words (about five to ten sentences). Each sentence gets a score between 0 and 1. The scores are averaged into the percentage teachers see.

A few details that matter to you:

- It only analyses **long-form prose**. Its 2023 FAQ says it does not assess bullet points, tables, poetry, code or short answers.
- On short documents of only a few hundred words, Turnitin says the prediction is "mostly all or nothing," because there is only one segment to judge.
- Scores from 1% to 19% are shown only as "%*" with no highlights, "to avoid potential incidences of false positives."
- In 2023 Turnitin said it only flags text when it is 98% confident, and that it may miss up to 15% of AI text as a result. It also said "real world use will differ from the lab tests."

What GPTZero says about its method

GPTZero made "perplexity" and "burstiness" famous. It now says it moved to a deep-learning model in autumn 2023, and that those two measures are now one of seven "indicators" it combines.

Perplexity and burstiness, in plain English

Perplexity is how surprised a language model is by each word. AI models write by choosing likely next words, so their text is usually low in surprise. GPTZero describes it as "how likely an AI model would have chosen the exact same set of words."

Burstiness is how much that surprise, and the sentence shape, changes across the text. People tend to write in bursts: a short sentence, then a long winding one, then a fragment. AI text is often smoother and more even.

EXAMPLE

Low surprise, even rhythm: "Time management is an important skill for students. It helps them balance their studies and personal life. Good time management can reduce stress and improve academic performance."

Higher surprise, uneven rhythm: "I learned time management the hard way. Three deadlines, one weekend, and a part-time shift at the pharmacy I had forgotten about. Now I plan Sundays like a train timetable, and I still miss things."

Both could be written by a person. The first one reads like a template, and that is the problem. Detectors cannot see who typed it. They only see how predictable it is.

A note on "humanizer" detection

Turnitin announced in August 2025 that its AI detection now includes "AI bypasser detection," trained on the patterns of leading humanizer tools. It supports English only for now.

What that means for you: running text through a rewriting tool is not a safe way to "pass." Detectors and rewriters keep changing. The durable approach is to write, or rewrite, in your own voice and keep a record of your process.

What the research says about accuracy

Study or statement	Finding
Liang et al., <i>Patterns</i> (2023)	Seven detectors misclassified over half of 91 TOEFL essays by non-native writers as AI (average false positive rate 61.22%). Essays by US 8th graders were judged almost perfectly.
Weber-Wulff et al., <i>International Journal for Educational Integrity</i> (2023)	Tested 14 tools, including Turnitin. Conclusion: "neither accurate nor reliable," with a main bias toward calling AI text human. Disguising AI text (for example by paraphrasing) made the tools perform much worse.
OpenAI (2023)	Withdrew its own classifier "due to its low rate of accuracy." On OpenAI's own test set it had caught 26% of AI text and wrongly flagged human text 9% of the time.
Turnitin (May 2023)	Document false positive rate under 1% for documents with over 20% AI writing; about 4% at the sentence level.
Turnitin (October 2023)	Its own study found no statistically significant difference in false positives between native and non-native writers for documents over 300 words. Shorter documents showed a larger gap.
Vanderbilt University (August 2023)	Turned off Turnitin's AI detector. It had submitted 75,000 papers to Turnitin in 2022; at a 1% false positive rate, around 750 papers could have been wrongly labeled.

PRO TIP

Notice the pattern. Tool makers report low error rates in their own tests. Independent researchers and universities report more problems. Both can be true: lab data is clean, real writing is messy. The safe conclusion is the one every side agrees on: a score is a reason to look closer, not proof.

Why non-native writers get flagged more

The Liang study is the clearest explanation. The essays that all seven detectors flagged had significantly **lower perplexity**: simpler, more common words in more predictable patterns. That is normal for someone writing carefully in a second language.

The researchers then ran two experiments:

- They used ChatGPT to make **US students' essays simpler**, "as if written by a non-native speaker." The average false positive rate jumped from 5.19% to 56.65%. Same ideas, simpler words, now "AI."
- They used ChatGPT to make the **TOEFL essays' vocabulary richer**. The false positive rate dropped from 61.22% to 11.77%.

The lesson is not "use AI to fix your English." The lesson is that **detectors partly measure vocabulary range and variety**, not honesty. If you write in a second language, you are not doing anything wrong. You are just more likely to be misread.

What helps, honestly:

- ☐ Write about **your** examples. Specific details are naturally less predictable than general statements.
- ☐ Use the words you would use when explaining the idea to a friend, not only the "safe" textbook phrases.
- ☐ Let your sentence length vary. A short sentence after a long one is fine in English.
- ☐ Do not over-polish. A tool that rewrites every sentence into perfect, formal English can make text more uniform, not less.
- ☐ Keep your drafts and version history. This is your best protection, whatever your English level.

The patterns that commonly trigger flags

Commercial detectors do not publish exactly what they look for. Our free detector at humanizeit.app does: it checks text against 40 patterns in five groups and shows which ones it found. These are the same things editors, teachers and readers notice when writing "feels like AI."

Two important points before the list:

- **One pattern means nothing.** Plenty of people say "furthermore." It is the combination, repeated across a text, that reads as machine-made.
- **Rhythm counts more than words.** In our detector, statistical signals (rhythm, repetition, variety) carry more weight than single words. Swapping "delve" for "explore" alone changes little.

1. Vocabulary

Pattern	Examples	Try instead
Words heavily overused by AI	delve, tapestry, pivotal, multifaceted, underscores, testament, cornerstone, embark, navigating, seamlessly	The plain word you would say out loud: "look at," "key," "shows."
Corporate buzzwords	leverage, streamline, robust, actionable, stakeholders, ecosystem, game-changing, cutting-edge	Say what actually happens: "cut the report time from 3 days to 1."
Promotional adjectives	nestled, breathtaking, vibrant, bustling, renowned, unparalleled	One concrete detail instead of two adjectives.
Melodramatic descriptors (creative writing)	ethereal, luminous, ephemeral, liminal, visceral, ineffable	Concrete images: what you saw, heard, smelled.

2. Phrases

Pattern	Examples	Try instead
Filler openers	"It's important to note that," "When it comes to," "In today's rapidly evolving..."	Delete them and start with the point.
"Plays a crucial role" family	"plays a crucial role," "one of the most important"	Say what role, and how you know.
Generic conclusions	"In conclusion," "The future is bright," "Only time will tell"	End on your strongest specific point or a real recommendation.

Pattern	Examples	Try instead
Hedging	"It could be argued that," "to some extent," "it depends on various factors"	Take a position, then name the real exception.
Vague attributions	"Experts say," "Studies show," "Research suggests"	Name the study, the author or the number.
The "despite challenges" formula	"Despite challenges, the industry continues to thrive."	Name the challenge and the evidence.
Assistant-style politeness	"Great question!", "I'd be happy to help"	Remove it. It never belongs in an essay or a post.
Stock casual openers	"Here's the thing," "Let's be honest," "Not gonna lie" (several in one text)	Casual is fine. Stock casual is not. Use your own phrasing.
Transition overuse	moreover, furthermore, additionally, consequently, in contrast	Let the logic carry the link, or use "and," "but," "so."

3. Structure

Pattern	What it looks like	Try instead
Repetitive sentence starters	Three or more sentences in a row starting with "The," "It" or "This."	Vary the openings. Start one with the subject, one with a time, one with a short clause.
Formulaic introduction and conclusion	Opens with "In today's digital age..."; closes by restating everything.	Open with a fact, a question you actually had, or a moment.
Uniform paragraphs	Every paragraph about the same length.	Let paragraph length follow the idea. One line is allowed.
Too many lists	Most of the text is bullet points.	Use lists for real lists, prose for reasoning.
Em dash overuse	Several dashes in a short text.	Use commas, full stops or brackets.
Colon overuse	Many sentences built as "Here is the key: ..."	Write the sentence straight.
Passive voice excess	"It was decided that mistakes were made."	Say who did what.
Groups of three	"fast, reliable, and secure" again and again.	Sometimes two items. Sometimes four. Sometimes one.
"Not just X, it's Y"	"It's not just a tool, it's a revolution."	Say what it is, once.

Pattern	What it looks like	Try instead
"Serves as" and "boasts"	"The library serves as a hub and boasts 3 floors."	"The library is a hub. It has 3 floors."
Empty "-ing" chains	"..., highlighting its importance and showcasing its impact."	Cut the tail, or give the actual result.
Rhetorical questions at paragraph edges	"But what does this mean for the future?"	Ask only questions you then really answer.

4. Meaning and tone

Pattern	What it looks like	Try instead
Over-explanation	Defining a term right after using it, for a reader who already knows it.	Write for your actual reader.
Always balanced	Every claim followed by "however" and the other side.	Show judgment. Say which side you find stronger and why.
Abstract language	Many abstract nouns ("growth," "innovation," "impact") and no examples.	Add a name, a number, a place, a date.
No personality	Long text with no question, no aside, no "I," no contractions.	Where your context allows it, let a bit of you show.
Too many qualifiers	"very," "really," "extremely," "incredibly"	Use a stronger noun or verb instead.

5. Statistics (the rhythm of the text)

Signal	What it measures	Why it matters
Low burstiness	Little variation in sentence length and shape.	The single strongest sign of "template" writing.
Uniform sentence length	Most sentences within a narrow range of word counts.	Real writing lurches: 4 words, then 30.
Average sentence length of 18 to 25 words	A typical AI comfort zone.	Not wrong on its own, suspicious when everything sits there.
Low word variety (type-token ratio)	The same words repeated across the text.	Shows limited vocabulary range, the effect seen in the Liang study.
Repeated three-word sequences	The same short phrases coming back.	Templates repeat; people usually rephrase.

Signal	What it measures	Why it matters
Predictable reading level	A middle-of-the-road readability score throughout.	Real texts rise and fall in difficulty.

What lowers the score in our detector: signs that only people usually leave. Contractions and informal markers where appropriate, a first-person voice, real variety in rhythm, and specific references such as asides in brackets. These are signals, not tricks. They appear naturally when you write about something you actually know.

How to write naturally

These habits make writing better for any reader. They also happen to make it less likely to be mistaken for a template.

Seven habits

1. **Start with your point.** No warm-up sentence about "today's world." Your first line should say something only your text says.
2. **Be specific.** Replace one general claim per paragraph with a fact: a name, a number, a date, a place, a quote.
3. **Use your experience.** Where your context allows, say what you saw, tried, got wrong or changed your mind about. No model has your memories.
4. **Vary your rhythm.** Read the text aloud. If every sentence sounds the same length, split one and join two others.
5. **Take a position.** Balance is good, endless balance is not. Say what you think, and why.
6. **Cut the filler.** Delete "It is important to note that" and see if the sentence still works. It almost always does.
7. **Edit by ear.** The final check is reading aloud. If you would never say it, rewrite it.

Before and after: a student paragraph

EXAMPLE

Before (template style): "In today's rapidly evolving world, social media plays a crucial role in the lives of young people. It is important to note that it has both positive and negative effects. On one hand, it fosters connection and creativity. On the other hand, it can lead to anxiety and decreased productivity. Ultimately, a balanced approach is essential."

After (same ideas, a real writer): "My younger brother spends about four hours a day on TikTok, by his own screen-time report. Some of it is good. He learned to edit video and now makes clips for his football club. But the night before his maths exam he was still scrolling at 1 a.m., and he knew it was a problem. I do not think the answer is a ban. I think it is a clear rule, like no phone after 11, that he agrees to himself."

What changed: a real example, two numbers, a contrast that comes from life, a clear opinion, mixed sentence lengths, no filler. Nothing was "disguised." It is simply better writing.

WATCH OUT

If your course, client or employer does not allow AI-written work, these habits are not a way around the rule. Write the work yourself, and if you used any tool, declare it. The free *False AI Flag Appeal Kit* explains how to keep evidence of your process.

Self-check before you submit

Run through this list on your final draft. It takes about ten minutes.

- ☐ My first sentence makes a real point. It does not start with "In today's..."
- ☐ Every paragraph has at least one specific detail (name, number, date, place or quote).
- ☐ I name my sources instead of "experts say" or "studies show."
- ☐ I have searched for "It is important to note," "plays a crucial role," "In conclusion" and removed most of them.
- ☐ No more than one or two of: delve, pivotal, multifaceted, tapestry, testament, seamless, leverage.
- ☐ Sentences vary in length. At least one is short. At least one is long.
- ☐ Not every paragraph is the same size.
- ☐ Three sentences in a row do not start with the same word.
- ☐ I do not list things in threes every time.
- ☐ I take a position somewhere, and say why.
- ☐ Where allowed, my own experience or perspective appears at least once.
- ☐ I read the whole text aloud, and changed anything I would never say.

- ☐ My drafts and version history are saved.
- ☐ I know the AI and tool policy that applies, and I followed it.
- ☐ If I used a tool (grammar checker, translator), I am ready to say so.

Honest limits

- **No tool can guarantee any outcome.** Not ours, not anyone's. Detectors are updated without notice and different detectors disagree with each other.
- **A low score is not a certificate.** It means one tool, on one day, did not find strong signals.
- **A high score is not proof.** It means one tool found patterns it associates with AI. Human review is needed, and the tool makers say so.
- **Short texts are less reliable.** Turnitin's own fairness study used a 300-word minimum, and Turnitin says short documents get "all or nothing" results.
- **Your process is your protection.** Version history, drafts and notes say far more about who wrote something than any score.

Check your own text free

Curious which patterns appear in your own writing?

- **Free AI detector:** humanizeit.app/ai-detector. Paste your text and see a breakdown of the patterns above, so you know which sentences read as formulaic and why. The instant scan runs in your browser, no signup.
- **Free humanizer:** humanizeit.app/free-ai-humanizer. Helpful for loosening stiff, over-formal wording into something closer to how you speak. Always review the result and keep your own facts and ideas.

Use them as a mirror, not a shield.

Sources

- Liang, W., Yuksekgonul, M., Mao, Y., Wu, E., Zou, J. (2023). "GPT detectors are biased against non-native English writers." *Patterns* 4(7). Preprint with the figures quoted here: <https://arxiv.org/abs/2304.02819>
- Weber-Wulff, D. et al. (2023). "Testing of detection tools for AI-generated text." *International Journal for Educational Integrity* 19: <https://doi.org/10.1007/s40979-023-00146-z> (preprint: <https://arxiv.org/abs/2306.15666>)

- OpenAI, "New AI classifier for indicating AI-written text" (31 Jan 2023, updated 20 Jul 2023): <https://openai.com/index/new-ai-classifier-for-indicating-ai-written-text/>
- Turnitin, "AI writing detection update from Turnitin's Chief Product Officer" (23 May 2023): <https://www.turnitin.co.uk/blog/ai-writing-detection-update-from-turnitins-chief-product-officer>
- Turnitin, "FAQs for Turnitin's AI writing detection capabilities" (2023): https://teaching.london.edu/wp-content/uploads/2023/03/AI_writing_FAQs_UK.pdf
- Turnitin Help Center, "Why is the AI Writing Detection report score showing as *%?" : <https://helpcenter.turnitin.com/hc/en-us/articles/46245332050317-Why-is-the-AI-Writing-Detection-report-score-showing-as>
- Turnitin, "New research: Turnitin's AI detector shows no statistically significant bias against English Language Learners" (26 Oct 2023): <https://www.turnitin.ca/blog/new-research-turnitin-s-ai-detector-shows-no-statistically-significant-bias-against-english-language-learners>
- Turnitin launches AI bypasser (humanizer) detection, Plagiarism Today (27 Aug 2025): <https://www.plagiarismtoday.com/2025/08/27/turnitin-launches-anti-ai-humanizer-feature/> and IT Brief (English-only support): <https://itbrief.co.uk/story/turnitin-updates-tool-to-detect-ai-bypassed-humanised-text>
- GPTZero, "Perplexity and burstiness: what is it?": <https://gptzero.me/news/perplexity-and-burstiness-what-is-it/>
- Google DeepMind, SynthID: <https://deepmind.google/technologies/synthid/> ; Axios, "Google DeepMind open sources its AI text watermarking tool" (24 Oct 2024): <https://www.axios.com/2024/10/24/google-deepmind-ai-text-watermarking-tool> ; TechCrunch on its limits with short, rewritten and translated text (23 Oct 2024): <https://techcrunch.com/2024/10/23/google-releases-tech-to-watermark-ai-generated-text/>
- Vanderbilt University, "Guidance on AI Detection and Why We're Disabling Turnitin's AI Detector" (16 Aug 2023): <https://www.vanderbilt.edu/brightspace/2023/08/16/guidance-on-ai-detection-and-why-were-disabling-turnitins-ai-detector/>
- Michigan State University, quoting Turnitin's disclaimer ("should not be used as the sole basis for adverse actions against a student"): <https://help.d2l.msu.edu/msu-docs/other-tools-at-msu/using-assignments-with-turnitin/ai-detection>
- Pattern list: Humanizeit's open detector, 40 patterns in five groups (vocabulary, phrase, structural, semantic, statistical), as shown in each result breakdown at <https://humanizeit.app/ai-detector>