

• FREE KIT

# False AI Flag Appeal Kit

---

What to do when your own writing is flagged as AI. A 24-hour plan, an evidence checklist, three email templates and a meeting script.

# Start here: the first 24 hours

You wrote it yourself. A detector says otherwise. Now a teacher, a client or an employer is asking questions.

This kit is for that moment. It works for students, freelancers and job seekers. It will not tell you how to "beat" a detector. It will help you do the one thing that actually settles these cases: **show how you wrote the work**, calmly and with evidence.

Two facts to hold on to:

- **A detector score is a guess, not a verdict.** Turnitin itself says its AI score "should not be used as the sole basis for adverse actions against a student" (sources at the end).
- **The person who flagged you usually wants to be fair.** Most of them are not trying to catch you. They saw a number and need help understanding it. Your job is to make that easy.

## WATCH OUT

Do not edit, rewrite or "clean up" the flagged document now. Do not run it through any rewriting tool. The original file, with its history, is your best evidence. Changing it can look like hiding something.

## The 24-hour plan

| When         | Do this                                                                                                                                                                                 | Why                                                                                                                                                |
|--------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------------------------------------------|
| Hour 0 to 2  | Read the message twice. Do not reply yet. Write down exactly what you are accused of, which tool was used, the score, and any deadline.                                                 | Replying while upset is the most common mistake. Facts first.                                                                                      |
| Hour 0 to 2  | Secure the evidence. Download a copy of the original file, screenshot the version history, save every related email. Leave the original where it is.                                    | A copied file does not bring its version history with it, older revisions can be merged, and some services keep only a limited number of versions. |
| Hour 2 to 8  | Find the rules. Look up the course syllabus, the AI policy, the academic integrity policy, or the client contract. Note what is allowed (spell checkers, Grammarly, translation tools). | You need to know which standard you are being held to.                                                                                             |
| Hour 2 to 8  | Build your evidence file using the checklist in this kit.                                                                                                                               | Evidence beats arguments.                                                                                                                          |
| Hour 8 to 24 | Send a short, polite reply using one of the templates. Ask for a meeting and for the specific passages that were flagged.                                                               | You show good faith and buy time to prepare properly.                                                                                              |

| When               | Do this                                                                                | Why                                                        |
|--------------------|----------------------------------------------------------------------------------------|------------------------------------------------------------|
| Before the meeting | Rehearse the meeting script out loud. Be ready to explain your argument without notes. | Being able to talk about your own work is strong evidence. |

## Questions to answer before you reply

- ☐ Which detector was used (Turnitin, GPTZero, Copyleaks, another)?
- ☐ What exactly was the score, and was it for the whole document or for specific sentences?
- ☐ Which passages were highlighted?
- ☐ Is this an informal question or a formal misconduct case?
- ☐ What is the deadline to respond, and who makes the decision?
- ☐ Which policy applies (course AI policy, university integrity policy, client contract, company rules)?
- ☐ Did you use any tool at all while writing (grammar checker, translator, AI for brainstorming)? Was it allowed?

### PRO TIP

If you did use a tool, such as Grammarly or a translator, say so plainly and early. Honesty about small things makes you believable about the big thing. Hiding it and having it come out later is far worse.

## Why detectors misfire

You do not need to become an expert. You need to understand enough to explain, calmly, why a score is not proof.

## How they work, in one page

AI detectors do not "find" AI. They **estimate a probability** that text looks like the text an AI model tends to produce.

- **Trained classifiers.** Most commercial tools are machine-learning models trained on examples of human and AI writing. Turnitin has described its process this way: it splits a submission into overlapping segments of a few hundred words, gives each sentence a score between 0 and 1, and averages them into an overall percentage.
- **Predictability ("perplexity").** AI models pick likely next words, so their text tends to be very predictable. GPTZero, one of the best known detectors, describes perplexity as a measure of "how likely an AI model would have chosen the exact same set of words."

- **Rhythm ("burstiness").** Humans tend to mix short and long sentences. AI text is often more even. GPTZero defines burstiness as how much writing patterns "vary over the entire document."

The problem: plenty of honest human writing is also predictable and even. Simple vocabulary, careful grammar, a fixed essay template, a formal tone. That is exactly how many students are taught to write, and how many people write in a second language.

## What the evidence says

| Source                                                                            | What it found or said                                                                                                                                                                                                                                                                                              |
|-----------------------------------------------------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Liang et al., Stanford, published in <i>Patterns</i> (2023)                       | Seven popular detectors were tested on 91 TOEFL essays written by non-native English speakers. They misclassified over half of them as AI-generated (average false positive rate 61.22%). 89 of the 91 essays (97.8%) were flagged by at least one detector. US 8th-grade essays were classified almost perfectly. |
| Same study                                                                        | When the authors used ChatGPT to make the US essays' word choices simpler, "as if written by a non-native speaker", the average false positive rate rose from 5.19% to 56.65%. Simpler wording alone made human writing look like AI.                                                                              |
| Turnitin, May 2023                                                                | Document false positive rate "less than 1%" for documents with over 20% AI writing; sentence-level false positive rate "~4%".                                                                                                                                                                                      |
| Turnitin help center (2026)                                                       | Scores between 1% and 19% are shown only as "%%" with no highlights, "to avoid potential incidences of false positives."                                                                                                                                                                                           |
| Turnitin guidance                                                                 | The score "should not be used as the sole basis for action" and Turnitin "does not make a determination of misconduct."                                                                                                                                                                                            |
| Turnitin's own FAQ                                                                | False positives can include "lists without a lot of structural variation, text that literally repeats itself, or text that has been paraphrased without developing new ideas." Short documents get "all or nothing" predictions.                                                                                   |
| Vanderbilt University, August 2023                                                | Disabled Turnitin's AI detector. Its reasoning: it submitted 75,000 papers to Turnitin in 2022, so a 1% false positive rate could have meant around 750 papers wrongly labeled. It also noted that detectors are more likely to flag writing by non-native English speakers.                                       |
| OpenAI, July 2023                                                                 | Withdrew its own AI text classifier "due to its low rate of accuracy." On its own test set it caught 26% of AI text and wrongly labeled human text as AI 9% of the time.                                                                                                                                           |
| Weber-Wulff et al., <i>International Journal for Educational Integrity</i> (2023) | Tested 14 detection tools, including Turnitin, and concluded they are "neither accurate nor reliable."                                                                                                                                                                                                             |

**PRO TIP**

Be fair when you use these facts. Turnitin published its own study (October 2023) reporting no statistically significant difference in false positives between native and non-native writers for documents over 300 words, and a larger gap for shorter documents. Detectors have been updated many times since 2023. Your point is not "detectors are useless." Your point is "a score alone cannot prove what I did or did not do, and here is what can."

## What "less than 1%" means for you

A 1% error rate sounds tiny. But it is 1 in 100. A large university or a busy hiring team processes thousands of documents. Some honest people will be flagged. That is simple arithmetic, and it is exactly why the tool makers say a human must review the evidence.

You might be one of those people. That is not a moral failing. It is a known limit of the technology.

# Build your evidence file

---

The strongest defense is **process evidence**: proof of how the work came to exist over time. A finished file shows the result. A history shows a human doing the work.

Collect everything first. Decide what to show later.

## The evidence checklist

### Process evidence (strongest)

- ☐ **Google Docs version history.** Open the document, then File, Version history, See version history (or click the "Last edit" link at the top). It lists edits with dates and times. You need edit access to the file to see it. Take screenshots of the timeline and of a few early versions. Do not restore or delete anything.
- ☐ **Microsoft Word version history.** Only available if the file is saved in OneDrive or SharePoint. Open the file, select the file name at the top, then Version history (older Office versions: File, History). Personal Microsoft accounts keep only the last 25 versions, so screenshot them now.
- ☐ **Earlier drafts and outlines.** Any separate files: "essay\_v1.docx", an outline, a rough plan.
- ☐ **File details.** Created and modified dates of the file and of earlier drafts.
- ☐ **Writing-process tools you already used.** For example, if Grammarly's Authorship feature was switched on while you wrote, it can produce a report and a replay of how the text was written. It only records while it is on, so it cannot help after the fact.
- ☐ **Emails or messages to yourself** with drafts attached, with their dates.

### Research trail

- ☐ Notes: handwritten pages (photograph them), notes apps, index cards.
- ☐ Sources you read: downloaded PDFs with your highlights or comments.
- ☐ Browser history for the research days (export or screenshot the relevant entries only).
- ☐ Library loans or database search history.

### People and feedback

- ☐ Writing center or tutor appointments, with dates.
- ☐ Feedback from a classmate, friend, colleague or editor on an earlier draft.
- ☐ Emails with your teacher or client about the topic while you were writing.

### Your voice over time

- ☐ Earlier work of yours in a similar style, ideally written in class, in an exam, or before AI tools were common.

- ☐ Other writing in the same register (emails, reports, blog posts) that shows how you normally write.

## Your knowledge

- ☐ A one-page summary of your argument, written from memory, so you can talk through it.
- ☐ A list of two or three choices you made and why (why this example, why this structure, which source you dropped).

### WATCH OUT

Only collect evidence that already exists. Never create a "draft" after the fact, change dates, or ask someone to say they saw you write it if they did not. Fabricated evidence turns a false accusation into a real offence.

## Evidence summary table

Make a one-page table like this. It is the document you hand over or attach.

| Evidence                                       | What it shows                                 | Date range      | Attached as |
|------------------------------------------------|-----------------------------------------------|-----------------|-------------|
| {evidence_1, e.g. Google Docs version history} | {e.g. 46 edit sessions over 9 days}           | {start_date} to | {file_name} |
| {evidence_2, e.g. outline and notes}           | {e.g. the structure existed before the draft} | {date}          | {file_name} |
| {evidence_3, e.g. tutor email}                 | {e.g. feedback on draft 2}                    | {date}          | {file_name} |

## If you have little or no history

It happens. You wrote in a desktop app without cloud saving, or on paper and typed it up at the end. Do not panic and do not invent anything.

- Lean on the other categories: notes, sources, feedback, earlier work.
- Offer to explain your work in a conversation, or to write a short piece on the topic under supervision.
- Ask what the decision will be based on, and point out that the tool makers say a score alone is not enough.

# Email templates

Short, calm, factual. Each template asks for the specific passages and offers evidence. Replace every {placeholder}. Delete any line that does not apply to you.

## Template 1: to an instructor or academic integrity office

### COPY-PASTE SCRIPT

**Subject:** {course\_code} {assignment\_name}: request to discuss the AI detection result

Dear {instructor\_name},

Thank you for letting me know about the AI detection result on my {assignment\_name}, submitted on {submission\_date}. I understand why the score raised a question, and I want to help resolve it.

I wrote this work myself. {If you used allowed tools, say so here, e.g. "I used Grammarly for spelling and grammar, which the course policy allows."} I have kept a record of my writing process, including {your\_evidence, e.g. the Google Docs version history from 2 to 11 March, my outline and research notes, and feedback from the writing center}.

Could you share which passages were flagged and which tool and score were used? I would also welcome a meeting, in person or online, where I can walk you through my drafts and explain my argument and sources.

I have attached a short summary of my evidence. I am happy to provide anything else that would help.

Kind regards,

{your\_full\_name}

{student\_id} · {program\_and\_year}

### PRO TIP

If this is already a formal case, send your reply to the person named in the letter and copy only who the procedure says to copy. Ask, politely, for a copy of the policy being applied and the timeline.

## Template 2: to a client or editor

### COPY-PASTE SCRIPT

**Subject:** {project\_name}: about the AI detection check

Hi {client\_name},

Thanks for flagging the detector result on {deliverable\_name}. I understand the concern, especially if your team or your own clients check content this way.

I wrote this piece myself. {optional\_tool\_note, e.g. "I used Grammarly for grammar checks, as we discussed."} I can share my working files to show how it was put together: {your\_evidence, e.g. the document's version history, my research notes and my interview notes}.

Detectors give a probability, not a confirmation. Detector makers such as Turnitin say their scores can be wrong and should not be used on their own. That said, I want you to be fully comfortable publishing this.

Two options: I can send the version history and notes today, or we can do a 15-minute call where I walk you through the draft. If any section reads too generic to you, I am also happy to revise it with more specific examples from {company\_or\_topic}.

Best,

## Template 3: to an employer or recruiter

Use this when a cover letter, a writing test or a work sample was flagged during hiring or at work.

#### **COPY-PASTE SCRIPT**

**Subject:** {position\_title} application: about my {cover\_letter / writing sample}

Dear {recruiter\_or\_manager\_name},

Thank you for your message about my {cover\_letter / writing sample} for the {position\_title} role. I appreciate the chance to respond.

I wrote it myself, drawing on my experience at {previous\_company\_or\_project}. {If relevant: "English is my second language, and I write carefully and formally, which I understand detectors sometimes read as AI."}

I would be glad to show my working drafts, or to complete a short writing task live or under time limits, on a topic of your choice. I would also be happy to talk through any part of the letter in an interview.

Thank you for considering this. I remain very interested in the role.

Best regards,  
{your\_full\_name}  
{phone} · {linkedin\_url}

#### **WATCH OUT**

Keep every message short and polite, even if you feel insulted. Do not attack the person or call the tool "a scam." Your reader is also the person who decides. Let your evidence do the arguing.

# What to say in the meeting

Meetings decide most of these cases. The goal is simple: by the end, the other person should see a real person who did real work.

## Before you walk in

- ☐ Bring your evidence summary table, printed or on screen.
- ☐ Have the version history open and ready to show.
- ☐ Re-read your work the night before. Know your main argument, your sources and why you chose them.
- ☐ Write down the three things you most want them to know.
- ☐ Check if you may bring someone (a student advisor, union rep, friend). Many procedures allow it. Your policy will say.

## The script

### COPY-PASTE SCRIPT

**Opening:** "Thank you for meeting with me. I understand why the score raised a question. I wrote this work myself, and I would like to show you how I did it."

**Your process (2 minutes):** "I started on {date} with {an outline / my notes from the reading}. Here is the version history. You can see I worked on it across {number} sessions. This section changed a lot after {feedback / rethinking my argument}."

**Your thinking:** "The main argument is {one sentence}. I used {source} because {reason}. I first tried {other idea} but dropped it because {reason}."

**Tools, honestly:** "I used {Grammarly / a translator / spell check} for {purpose}. I did not use AI to write the text." Or, if true: "I used {tool} to {brainstorm topics}, which I understood was allowed under {policy}."

**The tool's limits, briefly:** "I understand detectors give a probability. Turnitin's own guidance says the score should not be the only basis for a decision, and research has found that writers with simpler vocabulary, including many non-native speakers, are flagged more often. That is why I wanted to show you my process directly."

**Your offer:** "If it helps, I am happy to explain any paragraph, answer questions about the sources, or write a short piece on this topic here."

**Close:** "What else would be useful for you to see? And what are the next steps and timeline?"

## Good questions to ask

| Ask                                                                      | Why it helps                                                                                  |
|--------------------------------------------------------------------------|-----------------------------------------------------------------------------------------------|
| "Which passages were flagged, and what score did each get?"              | Lets you show the history of those exact passages.                                            |
| "Is the score the only concern, or is there something else in the work?" | Finds out if there is a second issue (sudden style change, a wrong citation) you can address. |
| "What evidence would resolve this for you?"                              | Turns the meeting into a checklist you can complete.                                          |
| "Which policy applies, and what are the next steps?"                     | Makes the process clear and fair for both sides.                                              |
| "Can I have a written summary of today's outcome?"                       | Protects you if the case goes further.                                                        |

## Avoid

- Saying "I never use any tools" if you used a grammar checker. Small untruths destroy trust.
- Long lectures about AI detection research. One sentence is enough. Then show your evidence.
- Getting defensive about style. If they say "this doesn't sound like you," show earlier work that does.
- Agreeing to anything you do not understand. You can say: "Can I take a day to think about that?"

## After the meeting

### COPY-PASTE SCRIPT

**Subject:** Thank you, and summary of our meeting on {date}

Dear {name},

Thank you for meeting with me today about {assignment\_or\_project}. To make sure I understood correctly: {one-line summary of what was agreed, e.g. "you will review the version history I shared and let me know your decision by {date}"}.

I have attached {the files we discussed}. Please let me know if anything else would help.

Kind regards,

## If the decision goes against you

- ☐ Ask for the decision and the reasons **in writing**.
- ☐ Find the appeal procedure and its deadline. Appeal deadlines can be short.

- ☐ Contact your student union, student advocate, ombudsperson or international student office. They have seen these cases before.
- ☐ In a job or client setting, ask whether a live writing task or a call would change the decision. If not, keep your evidence and move on. One recruiter's tool is not the whole market.

## Prevention habits for next time

A few habits make a future flag easy to resolve. None of them take much time.

- ☐ **Write in a tool that keeps history.** Google Docs, or Word saved to OneDrive. Write in one file from first note to final draft so the history tells the whole story. Personal OneDrive keeps only the last 25 versions, so also save a dated copy at each milestone (outline, first draft, final).
- ☐ **Keep a short writing log.** Date, time, what you did. One line per session is enough.
- ☐ **Save your notes and sources** in one folder per assignment or project.
- ☐ **Ask about the AI policy in writing** at the start. "Is Grammarly allowed? Is a translator allowed?" Keep the answer.
- ☐ **Turn on a process record if you want one.** Grammarly Authorship (available on its free plan) can track how text gets onto the page in Google Docs, Word or Grammarly's editor, but only while it is switched on.
- ☐ **Declare the tools you used.** Some courses ask for a short process note. Write one even if yours does not.
- ☐ **Keep earlier work.** It shows your voice over time.
- ☐ **Write in your own voice.** Use your own examples, your own experience and specific details. This is better writing anyway, and it is harder to mistake for a template. The free *AI Detection Field Guide 2026* covers this in detail.
- ☐ **Check before you submit.** Run your own text through a free detector that explains its result, so you know which passages read as formulaic. Then improve them by adding your own detail, not by disguising anything.

## A simple writing log

| Date   | Time          | What I did                                 | Where the evidence is |
|--------|---------------|--------------------------------------------|-----------------------|
| {date} | {start}-{end} | Read {source}, took notes on {topic}       | notes/{file}          |
| {date} | {start}-{end} | Outline, three main points                 | Google Doc v1         |
| {date} | {start}-{end} | Drafted intro and section 1                | Google Doc history    |
| {date} | {start}-{end} | Writing center feedback, revised section 2 | email from            |

# A note on honesty

---

This kit is for people who did the work and need to show it.

Present only genuine evidence. Never backdate files, create fake drafts, or ask anyone to vouch for something they did not see. If you used AI in a way your course, client or employer did not allow, the honest path is to say so and ask what can be done. Most institutions treat honesty far better than a cover-up that is later discovered.

If your rules allow AI help, use it openly and declare it. If they do not, don't use it for the writing itself.

## Check your own text free

---

Want to know how your writing might read to a detector before someone else checks it?

- **Free AI detector:** [humanizeit.app/ai-detector](https://humanizeit.app/ai-detector). It shows *why* text reads as AI, pattern by pattern, so you can see which sentences look formulaic. No signup.
- **Free humanizer:** [humanizeit.app/free-ai-humanizer](https://humanizeit.app/free-ai-humanizer). Useful when you want help turning stiff, formal wording into something that sounds more like you. Always check the result and keep your own ideas and facts.

No tool, ours included, can guarantee what another detector will say. Your best protection is your own voice plus a clear record of how you wrote.

## Sources

---

- Liang, W., Yuksekgonul, M., Mao, Y., Wu, E., Zou, J. (2023). "GPT detectors are biased against non-native English writers." *Patterns* 4(7). Preprint: <https://arxiv.org/abs/2304.02819>
- Turnitin, "AI writing detection update from Turnitin's Chief Product Officer" (Annie Chechitelli, 23 May 2023): <https://www.turnitin.co.uk/blog/ai-writing-detection-update-from-turnitins-chief-product-officer>
- Turnitin, "FAQs for Turnitin's AI writing detection capabilities" (2023, hosted by London Business School): [https://teaching.london.edu/wp-content/uploads/2023/03/AI\\_writing\\_FAQs\\_UK.pdf](https://teaching.london.edu/wp-content/uploads/2023/03/AI_writing_FAQs_UK.pdf)
- Turnitin Help Center, "Why is the AI Writing Detection report score showing as \*%?" : <https://helpcenter.turnitin.com/hc/en-us/articles/46245332050317-Why-is-the-AI-Writing-Detection-report-score-showing-as>
- Turnitin guidance quoted by Michigan State University: "should not be used as the sole basis for adverse actions against a student": <https://help.d2l.msu.edu/msu-docs/other-tools-at-msu/using-assignments-with-turnitin/ai-detection>

- Turnitin, "New research: Turnitin's AI detector shows no statistically significant bias against English Language Learners" (26 Oct 2023): <https://www.turnitin.ca/blog/new-research-turnitin-s-ai-detector-shows-no-statistically-significant-bias-against-english-language-learners>
- Vanderbilt University, "Guidance on AI Detection and Why We're Disabling Turnitin's AI Detector" (16 Aug 2023): <https://www.vanderbilt.edu/brightspace/2023/08/16/guidance-on-ai-detection-and-why-were-disabling-turnitins-ai-detector/>
- OpenAI, "New AI classifier for indicating AI-written text" (31 Jan 2023, updated 20 Jul 2023): <https://openai.com/index/new-ai-classifier-for-indicating-ai-written-text/>
- Weber-Wulff, D. et al. (2023). "Testing of detection tools for AI-generated text." *International Journal for Educational Integrity* 19: <https://doi.org/10.1007/s40979-023-00146-z> (preprint: <https://arxiv.org/abs/2306.15666>)
- GPTZero, "Perplexity and burstiness: what is it?": <https://gptzero.me/news/perplexity-and-burstiness-what-is-it/>
- MIT Sloan Teaching & Learning Technologies, "AI Detectors Don't Work. Here's What to Do Instead.": <https://mitsloanedtech.mit.edu/ai/teach/ai-detectors-dont-work/>
- Google Docs Help, "Find what's changed in a file": <https://support.google.com/docs/answer/190843>
- Microsoft Support, "View previous versions of Office files": <https://support.microsoft.com/topic/5c1e076f-a9c9-41b8-8ace-f77b9642e2c2>
- Grammarly Authorship: <https://www.grammarly.com/authorship>

*This kit is general guidance, not legal advice. Procedures differ between institutions, employers and countries. Always follow the official process that applies to you.*